

EXP NO: 5

INSTALLATION OF HIVE

\$wget <https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz>

```
Activities Terminal Sep 18 11:22
hadoop@krish-VirtualBox: ~
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

hadoop@krish-VirtualBox: ~$ cd
hadoop@krish-VirtualBox: ~$ wget https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz
--2024-09-18 10:48:58-- https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz
Resolving archive.apache.org (archive.apache.org)... 65.108.204.189, 2a01:4f9:1a:a084::2
Connecting to archive.apache.org (archive.apache.org)|65.108.204.189|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 278813748 (266M) [application/x-gzip]
Saving to: 'apache-hive-3.1.2-bin.tar.gz'

apache-hive-3.1.2-bin.tar.gz 100%[=====] 265.90M 155KB/s in 29
n 48s

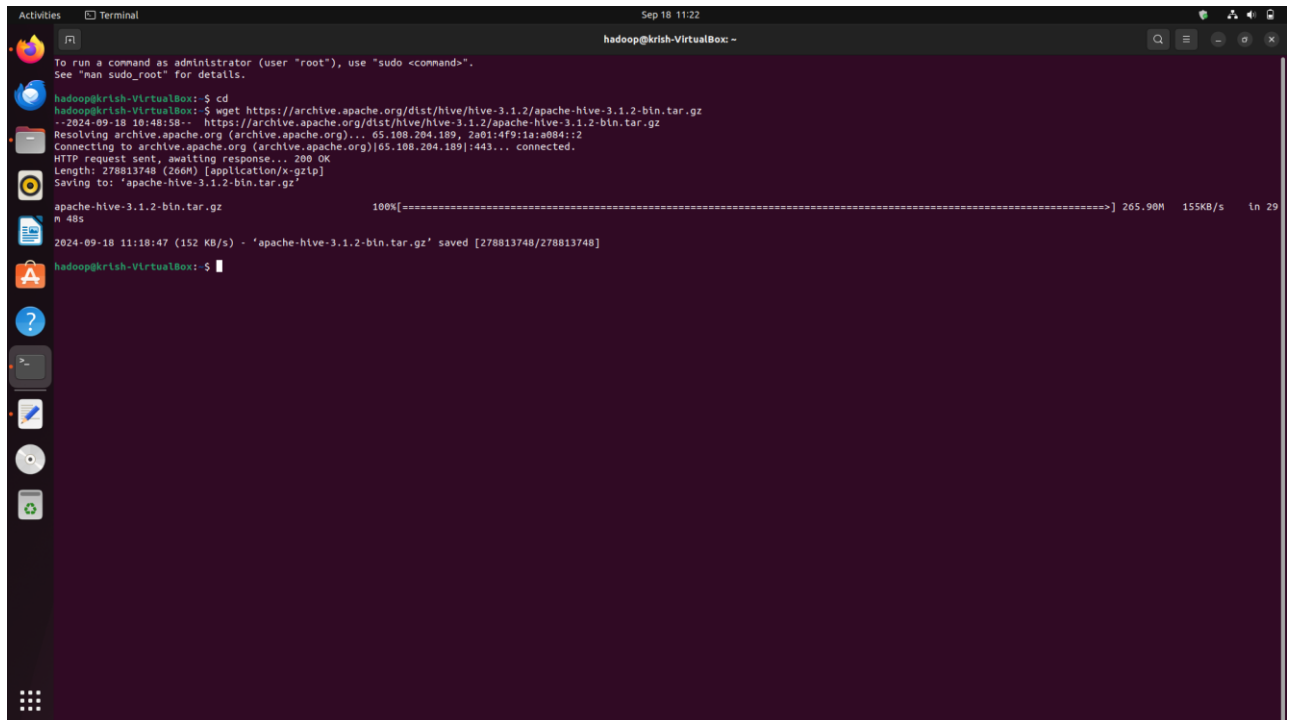
2024-09-18 11:10:47 (152 KB/s) - 'apache-hive-3.1.2-bin.tar.gz' saved [278813748/278813748]

hadoop@krish-VirtualBox: ~$
```

\$tar xvfz apache-hive-3.1.2-bin.tar.gz

```
Activities Terminal Sep 18 12:03
hadoop@krish-VirtualBox: ~
hadoop@krish-VirtualBox: ~$ tar xvfz apache-hive-3.1.2-bin.tar.gz
apache-hive-3.1.2-bin/LICENSE
apache-hive-3.1.2-bin/NOTICE
apache-hive-3.1.2-bin/RELEASE_NOTES.txt
apache-hive-3.1.2-bin/binary-package/licenses/asn-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/com.google.protobuf-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/com.ibm.icu.icu4j-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/com.sun.jersey-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/com.thoughtworks.paranamer-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/javaw.transaction.transaction-api-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/javalution-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/jline-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/NOTICE
apache-hive-3.1.2-bin/binary-package/licenses/org.abego.treelayout.core-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/org.apache.antlr4-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/org.apache.antlr4.stringtemplate-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/org.codehaus.janino-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/org.jamon.jamon-runtime-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/org.jruby-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/org.mozilla.rhino-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/org.slf4j-LICENSE
apache-hive-3.1.2-bin/binary-package/licenses/sqlite-LICENSE
apache-hive-3.1.2-bin/examples/files/2000_cols_data.csv
apache-hive-3.1.2-bin/examples/files/3col_data.txt
apache-hive-3.1.2-bin/examples/files/4col_data.txt
apache-hive-3.1.2-bin/examples/files/agg_01-p1.txt
apache-hive-3.1.2-bin/examples/files/agg_01-p2.txt
apache-hive-3.1.2-bin/examples/files/agg_01-p3.txt
apache-hive-3.1.2-bin/examples/files/alltypes.txt
apache-hive-3.1.2-bin/examples/files/alltypes2.txt
apache-hive-3.1.2-bin/examples/files/apache.access.2.log
apache-hive-3.1.2-bin/examples/files/apache.access.log
apache-hive-3.1.2-bin/examples/files/archive_corrupt.rc
apache-hive-3.1.2-bin/examples/files/array_table.txt
apache-hive-3.1.2-bin/examples/files/avro_charvarchar.txt
apache-hive-3.1.2-bin/examples/files/avro_date.txt
apache-hive-3.1.2-bin/examples/files/avro_historical_timestamp_legacy.avro
apache-hive-3.1.2-bin/examples/files/avro_historical_timestamp_new.avro
apache-hive-3.1.2-bin/examples/files/avro_timestamp.txt
apache-hive-3.1.2-bin/examples/files/AvroPrintableInList.parquet
apache-hive-3.1.2-bin/examples/files/AvroSingleFieldGroupInList.parquet
apache-hive-3.1.2-bin/examples/files/binary.txt
apache-hive-3.1.2-bin/examples/files/bool.txt
apache-hive-3.1.2-bin/examples/files/bool_literal.txt
apache-hive-3.1.2-bin/examples/files/cbo_t1.txt
apache-hive-3.1.2-bin/examples/files/cbo_t2.txt
apache-hive-3.1.2-bin/examples/files/cbo_t3.txt
apache-hive-3.1.2-bin/examples/files/cbo_t4.txt
apache-hive-3.1.2-bin/examples/files/cbo_t5.txt
apache-hive-3.1.2-bin/examples/files/cbo_t6.txt
apache-hive-3.1.2-bin/examples/files/char_varchar_udf.txt
apache-hive-3.1.2-bin/examples/files/complex.seq
apache-hive-3.1.2-bin/examples/files/covar_tab.txt
apache-hive-3.1.2-bin/examples/files/create_nested_type.txt
```

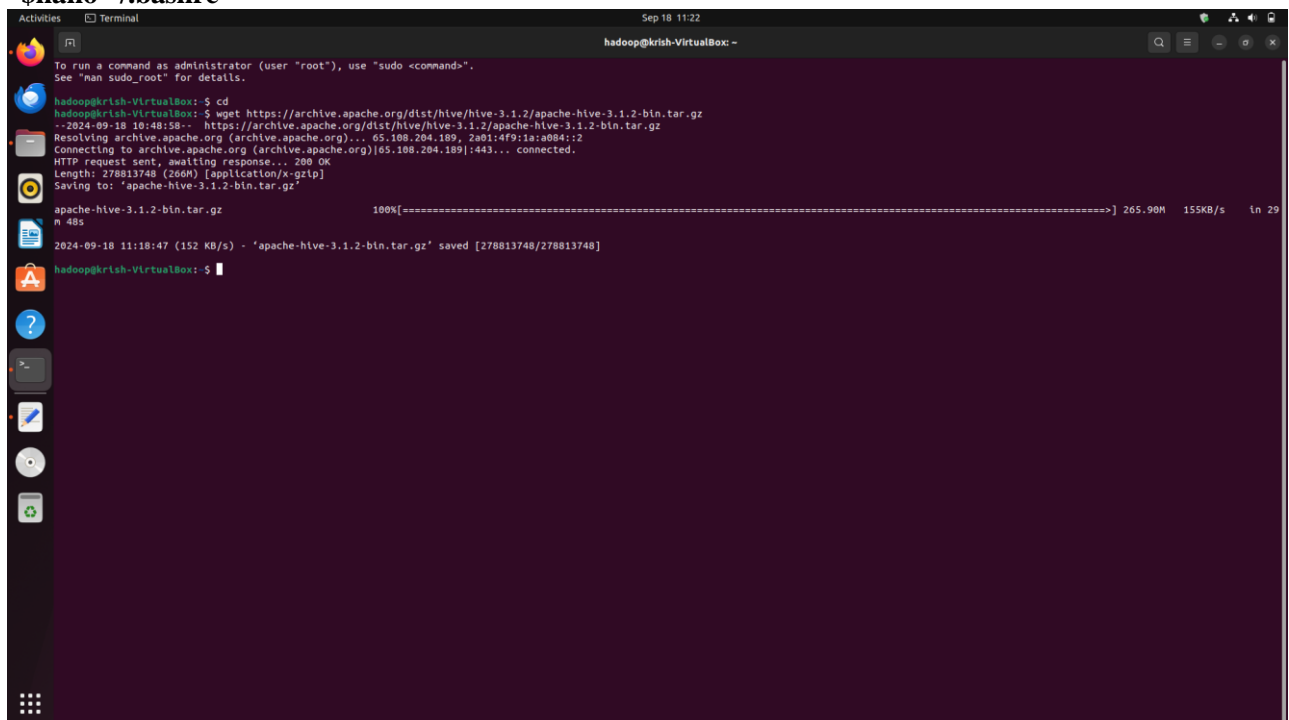
\$mv apache-hive-3.1.2-bin hive



A terminal window titled 'hadoop@krish-VirtualBox: ~' showing the process of downloading Apache Hive 3.1.2 binaries. The user runs 'cd' and then 'wget https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz'. The terminal shows the file being resolved, connected to, and saved. A progress bar indicates 100% completion. The file size is 265.90M and the download speed is 155KB/s. The terminal also shows the file being saved to 'apache-hive-3.1.2-bin.tar.gz'.

```
hadoop@krish-VirtualBox: ~  
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.  
hadoop@krish-VirtualBox: ~$ cd  
hadoop@krish-VirtualBox: ~$ wget https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz  
--2024-09-18 10:48:58-- https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz  
Resolving archive.apache.org (archive.apache.org)... 65.108.204.189, 2a01:4f9:1a:a084::2  
Connecting to archive.apache.org (archive.apache.org)[65.108.204.189]:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 278813748 (266M) [application/x-gzip]  
Saving to: 'apache-hive-3.1.2-bin.tar.gz'  
  
apache-hive-3.1.2-bin.tar.gz      100%[=====] 265.90M  155KB/s   in 29  
n 48s  
2024-09-18 11:18:47 (152 KB/s) - 'apache-hive-3.1.2-bin.tar.gz' saved [278813748/278813748]  
hadoop@krish-VirtualBox: ~$
```

\$nano ~/.bashrc



A terminal window titled 'hadoop@krish-VirtualBox: ~' showing the nano editor editing the ~/.bashrc file. The terminal shows the same download process as the previous screenshot, but the user is now in the nano editor. The terminal also shows the file being saved to 'apache-hive-3.1.2-bin.tar.gz'.

```
hadoop@krish-VirtualBox: ~  
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.  
hadoop@krish-VirtualBox: ~$ cd  
hadoop@krish-VirtualBox: ~$ wget https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz  
--2024-09-18 10:48:58-- https://archive.apache.org/dist/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz  
Resolving archive.apache.org (archive.apache.org)... 65.108.204.189, 2a01:4f9:1a:a084::2  
Connecting to archive.apache.org (archive.apache.org)[65.108.204.189]:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 278813748 (266M) [application/x-gzip]  
Saving to: 'apache-hive-3.1.2-bin.tar.gz'  
  
apache-hive-3.1.2-bin.tar.gz      100%[=====] 265.90M  155KB/s   in 29  
n 48s  
2024-09-18 11:18:47 (152 KB/s) - 'apache-hive-3.1.2-bin.tar.gz' saved [278813748/278813748]  
hadoop@krish-VirtualBox: ~$
```

\$source ~/.bashrc

\$cd hive

\$cp conf/hive-default.xml.template conf/hive-site.xml

\$nano conf/hive-site.xml

\$start-all.sh

\$ jps

```

krish@krish-VirtualBox: ~
Starting namenodes on [localhost]
Starting datanodes
localhost: datanode is running as process 4172. Stop it first and ensure /tmp/hadoop-krish-datanode.pid file is empty before retry.
Starting secondary namenodes [krish-VirtualBox]
krish-VirtualBox: secondarynamenode is running as process 4366. Stop it first and ensure /tmp/hadoop-krish-secondarynamenode.pid file is empty before retry.
Starting resourcemanager
resourcemanager is running as process 4573. Stop it first and ensure /tmp/hadoop-krish-resourcemanager.pid file is empty before retry.
Starting nodemanagers
localhost: nodemanager is running as process 4715. Stop it first and ensure /tmp/hadoop-krish-nodemanager.pid file is empty before retry.
krish@krish-VirtualBox:~$ stop-all.sh
WARNING: Stopping all Apache Hadoop daemons as krish in 10 seconds.
WARNING: Use CTRL-C to abort.
Stopping namenodes on [localhost]
Stopping datanodes
Stopping secondary namenodes [krish-VirtualBox]
Stopping nodemanagers
Stopping resourcemanager
krish@krish-VirtualBox:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as krish in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [krish-VirtualBox]
Starting resourcemanager
Starting nodemanagers
krish@krish-VirtualBox:~$ jps
8721 Jps
8009 SecondaryNameNode
7626 NameNode
7773 DataNode
8351 NodeManager
8207 ResourceManager
  
```

\$hdfs dfs -mkdir /user

\$hdfs dfs -mkdir /user/hive

\$hdfs dfs -mkdir /user/hive/warehouse

\$hdfs dfs -mkdir /user/tmp

\$hdfs dfs -chmod g+w /user/tmp

\$hdfs dfs -chmod g+w /user/hive/warehouse

\$cd \$HIVE_HOME/bin

\$schematool -initSchema -dbType derby

```

hadoop@krish-VirtualBox: ~/hive/bin
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/jaxb-impl-2.2.3-1.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/jackson-jaxrs-1.9.2.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/jackson-xc-1.9.2.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/jersey-core-1.19.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/jersey-servlet-1.19.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/hive-webhcat-3.1.2.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/wadl-resourcecdoc-doctet-1.4.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/xercesImpl-2.9.1.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/xml-apis-1.3.04.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/commons-exec-1.1.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/svr/lib/jul-to-slf4j-1.7.10.jar
apache-hive-3.1.2-bin/hcatalog/share/webhcat/java-client/hive-webhcat-java-client-3.1.2.jar
hadoop@krish-VirtualBox:~$ mv apache-hive-3.1.2 hive
mv: cannot stat 'apache-hive-3.1.2': No such file or directory
hadoop@krish-VirtualBox:~$ ls
Desktop  Downloads  hadoopdata  Music  pig  pig-0.16.0.tar.gz  pig-0.17.0.tar.gz  Public  sample.txt  Templates
hadoop@krish-VirtualBox:~$ cd hive
hadoop@krish-VirtualBox:~/hive$ cp conf/hive-default.xml.template conf/hive-site.xml
hadoop@krish-VirtualBox:~/hive$ nano conf/hive-site.xml
hadoop@krish-VirtualBox:~/hive$ hdfs dfs -mkdir /user
Refused: Call From krish-VirtualBox/127.0.1.1 to localhost:9000 failed on connection exception: java.net.ConnectException: Connection refused; For more details see: http://wiki.apache.org/hadoop/Connection
hadoop@krish-VirtualBox:~/hive$ jps
2497 Jps
hadoop@krish-VirtualBox:~/hive$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [krish-VirtualBox]
Starting resourcemanager
Starting nodemanagers
hadoop@krish-VirtualBox:~/hive$ hdfs dfs -mkdir /user
hadoop@krish-VirtualBox:~/hive$ hdfs dfs -mkdir /user/hive
hadoop@krish-VirtualBox:~/hive$ hdfs dfs -mkdir /user/hive/warehouse
hadoop@krish-VirtualBox:~/hive$ hdfs dfs -mkdir /user/tmp
hadoop@krish-VirtualBox:~/hive$ hdfs dfs -chmod g+w /user/tmp
hadoop@krish-VirtualBox:~/hive$ hdfs dfs -chmod g+w /user/hive/warehouse
hadoop@krish-VirtualBox:~/hive$ cd $HIVE_HOME/bin
hadoop@krish-VirtualBox:~/hive$ schematool -initSchema -dbtype derby
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Metastore connection URL: jdbc:derby:databaseName=metastore_db;create=true
Metastore connection Driver: org.apache.derby.jdbc.EmbeddedDriver
Metastore connection User: APP
  
```

\$hive

```

Activities  Terminal  Sep 18 12:20
hadoop@krish-VirtualBox: ~/hive/bin

Initialization script completed
schemaTool completed
hadoop@krish-VirtualBox:~/hive/bin$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = f9bdf2e7-d022-4b43-8660-4967ac66a2ab

Logging initialized using configuration in jar:file:/home/hadoop/hive/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive Session ID = d644c7a9-bb2b-4a14-9c70-3758ec82bb35
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> show databases;
OK
default
Time taken: 1.105 seconds, Fetched: 1 row(s)
hive>

```

>show databases;

```

Activities  Terminal  Sep 18 12:20
hadoop@krish-VirtualBox: ~/hive/bin

Initialization script completed
schemaTool completed
hadoop@krish-VirtualBox:~/hive/bin$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = f9bdf2e7-d022-4b43-8660-4967ac66a2ab

Logging initialized using configuration in jar:file:/home/hadoop/hive/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive Session ID = d644c7a9-bb2b-4a14-9c70-3758ec82bb35
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> show databases;
OK
default
Time taken: 1.105 seconds, Fetched: 1 row(s)
hive>

```

> **CREATE DATABASE financials;**

> **use financials;**

> **CREATE TABLE finance_table(id INT, name STRING);**

> **INSERT INTO finance_table VALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');**

```

Activities Terminal Sep 18 14:15
hadoop@krish-VirtualBox: ~/hive/bin

Initialization script completed
schemaTool completed
hadoop@krish-VirtualBox:~/hive/bin$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/hive/lib/hive-log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 6af165f1-39bd-4b22-9958-f4477b2557ed

Logging initialized using configuration in jar:file:/home/hadoop/hive/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Hive Session ID = d99da8f9-0b51-4f4f-b8a1-ads2c2c69d508
hive> show databases;
OK
default
Time taken: 1.191 seconds, Fetched: 1 row(s)
hive> CREATE DATABASE financials;
OK
Time taken: 0.261 seconds
hive> use financials;
OK
Time taken: 0.09 seconds
hive> CREATE TABLE finance_table( id INT, name STRING );
OK
Time taken: 1.305 seconds
hive> INSERT INTO finance_table VALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');
Query ID = hadoop_20240918141355_7cc136e1-6955-4787-b44f-ea060434980a
Total Jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reducers=<number>
Job running in-process (local Hadoop)
2024-09-18 14:14:03,855 Stage-1 map = 0%, reduce = 0%
2024-09-18 14:14:07,280 Stage-1 map = 100%, reduce = 0%
2024-09-18 14:14:08,323 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local116690273_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://localhost:9000/user/hive/warehouse/financials.db/finance_table/.hive-staging_hive_2024-09-18_14-13-55_239_3003215123774537445-1/-ext-10000
Loading data to table financials.finance_table
MapReduce Jobs Launched:
Stage-Stage-1: HDFS Read: 0 HDFS Write: 208 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 15.305 seconds
hive>

```

> **CREATE VIEW myview AS SELECT name, id FROM finance_table;**

> **SELECT * FROM myview;**

```

Activities Terminal Sep 18 14:16
hadoop@krish-VirtualBox: ~/hive/bin

SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 6af165f1-39bd-4b22-9958-f4477b2557ed

Logging initialized using configuration in jar:file:/home/hadoop/hive/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Hive Session ID = d99da8f9-0b51-4f4f-b8a1-ads2c2c69d508
hive> show databases;
OK
default
Time taken: 1.191 seconds, Fetched: 1 row(s)
hive> CREATE DATABASE financials;
OK
Time taken: 0.261 seconds
hive> use financials;
OK
Time taken: 0.09 seconds
hive> CREATE TABLE finance_table( id INT, name STRING );
OK
Time taken: 1.305 seconds
hive> INSERT INTO finance_table VALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');
Query ID = hadoop_20240918141355_7cc136e1-6955-4787-b44f-ea060434980a
Total Jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reducers=<number>
Job running in-process (local Hadoop)
2024-09-18 14:14:03,855 Stage-1 map = 0%, reduce = 0%
2024-09-18 14:14:07,280 Stage-1 map = 100%, reduce = 0%
2024-09-18 14:14:08,323 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local116690273_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://localhost:9000/user/hive/warehouse/financials.db/finance_table/.hive-staging_hive_2024-09-18_14-13-55_239_3003215123774537445-1/-ext-10000
Loading data to table financials.finance_table
MapReduce Jobs Launched:
Stage-Stage-1: HDFS Read: 0 HDFS Write: 208 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 15.305 seconds
hive> CREATE VIEW myview AS SELECT name, id FROM finance_table;
OK
Time taken: 0.427 seconds
hive> SELECT * FROM myview;
OK
Alice 1
Bob 2
Charlie 3
Time taken: 0.403 seconds, Fetched: 3 row(s)
hive>

```

> **DESCRIBE finance_table;**

> **ALTER TABLE finance_table ADD COLUMNS (age INT);**

```

Activities Terminal Sep 18 14:17
hadoop@krish-VirtualBox: ~/hive/bin

default
Time taken: 1.191 seconds, Fetched: 1 row(s)
hive> CREATE DATABASE financials;
OK
Time taken: 0.261 seconds
hive> use financials;
OK
Time taken: 0.09 seconds
hive> CREATE TABLE finance_table( id INT, name STRING );
OK
Time taken: 1.305 seconds
hive> INSERT INTO finance_table VALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');
Query ID = hadoop_20240918141355_7cc136e1-6955-4707-b44f-ea060434980a
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2024-09-18 14:14:03,855 Stage-1 map = 0%, reduce = 0%
2024-09-18 14:14:07,280 Stage-1 map = 100%, reduce = 0%
2024-09-18 14:14:08,323 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local116690273_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://localhost:9000/user/hive/warehouse/financials.db/finance_table/.hive-staging_hive_2024-09-18-14-13-55_239_3003215123774537445-1/-ext-10000
Loading data to table financials.finance_table
MapReduce Jobs Launched:
Stage-Stage-1:  HDFS Read: 0 HDFS Write: 208 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 15.305 seconds
hive> CREATE VIEW myview AS SELECT name, id FROM finance_table;
OK
Time taken: 0.427 seconds
hive> SELECT * FROM myview;
OK
Alice      1
Bob        2
Charlie    3
Time taken: 0.403 seconds, Fetched: 3 row(s)
hive> DESCRIBE finance_table;
OK
id          int
name        string
Time taken: 0.049 seconds, Fetched: 2 row(s)
hive> ALTER TABLE finance_table ADD COLUMN (age INT);
OK
Time taken: 0.185 seconds
hive>

```

>quit;

```

Activities Terminal Sep 18 14:17
hadoop@krish-VirtualBox: ~/hive/bin

default
Time taken: 1.191 seconds, Fetched: 1 row(s)
hive> CREATE DATABASE financials;
OK
Time taken: 0.261 seconds
hive> use financials;
OK
Time taken: 0.09 seconds
hive> CREATE TABLE finance_table( id INT, name STRING );
OK
Time taken: 1.305 seconds
hive> INSERT INTO finance_table VALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');
Query ID = hadoop_20240918141355_7cc136e1-6955-4707-b44f-ea060434980a
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2024-09-18 14:14:03,855 Stage-1 map = 0%, reduce = 0%
2024-09-18 14:14:07,280 Stage-1 map = 100%, reduce = 0%
2024-09-18 14:14:08,323 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local116690273_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://localhost:9000/user/hive/warehouse/financials.db/finance_table/.hive-staging_hive_2024-09-18-14-13-55_239_3003215123774537445-1/-ext-10000
Loading data to table financials.finance_table
MapReduce Jobs Launched:
Stage-Stage-1:  HDFS Read: 0 HDFS Write: 208 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 15.305 seconds
hive> CREATE VIEW myview AS SELECT name, id FROM finance_table;
OK
Time taken: 0.427 seconds
hive> SELECT * FROM myview;
OK
Alice      1
Bob        2
Charlie    3
Time taken: 0.403 seconds, Fetched: 3 row(s)
hive> DESCRIBE finance_table;
OK
id          int
name        string
Time taken: 0.049 seconds, Fetched: 2 row(s)
hive> ALTER TABLE finance_table ADD COLUMN (age INT);
OK
Time taken: 0.185 seconds
hive>

```